



JSPM's
RAJARSHI SHAHU COLLEGE OF ENGINEERING
TATHAWADE, PUNE-33



DEPARTMENT OF COMPUTER ENGINEERING

Case Study

On

Email Spam Detection

-Submitted by-

Name of Student: Siddharth Vikram Maral

Roll No: CS3172

PRN No: RBTL23CS152

Class and Division: TY - A

Course: Machine Learning

1. Title

Email Spam Detection

2. Background/ Introduction (about details of application)

Email spam detection is an essential application of machine learning that focuses on distinguishing legitimate emails from unwanted or harmful ones. With the exponential growth of digital communication, spam emails have become a major threat, often carrying phishing links, malware, or fraudulent content. Traditional rule-based filters are no longer sufficient to handle the complexity and volume of modern spam messages. Machine learning techniques provide a more intelligent and adaptive solution by learning from past email patterns, analyzing textual content, sender behavior, and message features to automatically detect spam. This not only enhances cybersecurity but also reduces user distraction and improves overall email management efficiency.

3. Problem Statement

The problem is to automatically identify and filter spam emails using machine learning techniques. The goal is to build an intelligent system that accurately classifies emails as spam or ham to enhance security and reduce unwanted messages.

4. Objectives

- To develop a machine learning model that accurately classifies emails as spam or ham.
- To preprocess and analyze email text data using natural language processing techniques.
- To evaluate multiple classifiers (Naive Bayes, SVM, Logistic Regression, etc.) and select the best-performing model.
- To improve email security and user productivity by automatically filtering unwanted messages.

5. Libraries required

- pandas
- numpy
- sklearn
- nltk
- matplotlib
- flask

6. Implementation

```
# 1 Email Spam Detection using Machine Learning
import os
import argparse
import glob
import re
import joblib
from pathlib import Path
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.pipeline import Pipeline
from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import LogisticRegression
from sklearn.svm import LinearSVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score,
confusion_matrix, classification_report
import matplotlib.pyplot as plt

# Optional: NLTK for stopwords/tokenization
try:
    import nltk
    from nltk.corpus import stopwords
except Exception:
    nltk = None

def normalize_label(x):
    if isinstance(x, (int, float)):
        return 'spam' if int(x) == 1 else 'ham'
    s = str(x).strip().lower()
    if s in ['spam', '1', 'true', 't', 's']:
        return 'spam'
    return 'ham'

def load_from_csv(csv_path: str) -> pd.DataFrame:
    df = pd.read_csv(csv_path)

    # Flexible column naming
    if 'text' not in df.columns:
        if 'message' in df.columns:
            df.rename(columns={'message': 'text'}, inplace=True)
        elif 'body' in df.columns:
            df.rename(columns={'body': 'text'}, inplace=True)
        else:
```

```

        raise ValueError("CSV must contain a 'text' column")

    if 'label' not in df.columns:
        for alt in ['target', 'class']:
            if alt in df.columns:
                df.rename(columns={alt: 'label'}, inplace=True)
                break
        else:
            raise ValueError("CSV must contain a 'label' column")

    df = df[['text', 'label']].dropna()
    df['label'] = df['label'].apply(lambda x: normalize_label(x))
    return df

def load_from_folders(base_path: str) -> pd.DataFrame:
    records = []
    for label in ['spam', 'ham']:
        folder = Path(base_path) / label
        if not folder.exists():
            continue

        for filepath in folder.glob('**/*.txt'):
            try:
                text = filepath.read_text(encoding='utf-8', errors='ignore')
            except Exception:
                text = ""
            records.append({'text': text, 'label': label})

    if not records:
        raise ValueError('No data found. Expecting ./data/spam/ and ./data/ham/ with text files')

    return pd.DataFrame(records)

def basic_clean(text: str) -> str:
    if not isinstance(text, str):
        return ""
    text = re.sub(r'<[^>]+>', '', text) # remove HTML tags
    text = re.sub(r'http\S+|www\.[^\s]+', '', text) # remove URLs
    text = re.sub(r'\S+@\S+', '', text) # remove emails
    text = re.sub(r'^0-9a-zA-Z\s', '', text) # remove special chars
    text = re.sub(r'\s+', ' ', text) # collapse spaces
    return text.strip().lower()

def prepare_data(data_dir: str = './data') -> pd.DataFrame:
    csv_path = os.path.join(data_dir, 'emails.csv')
    if os.path.exists(csv_path):
        print(f'Loading dataset from {csv_path}')

```

```

    df = load_from_csv(csv_path)
else:
    print(f'CSV not found. Trying folder load from {data_dir}')
    df = load_from_folders(data_dir)

print(f'Loaded {len(df)} records')
df['text_clean'] = df['text'].apply(basic_clean)
return df

def build_vectorizer(max_features: int = 20000):
    vect = TfidfVectorizer(max_features=max_features, ngram_range=(1, 2), stop_words='english')
    return vect

def evaluate_model(model, X_test, y_test, model_name: str, plot_prefix: str = 'plot'):
    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred, pos_label='spam')
    rec = recall_score(y_test, y_pred, pos_label='spam')
    f1 = f1_score(y_test, y_pred, pos_label='spam')

    print(f"\n== {model_name} ==")
    print(f"Accuracy: {acc:.4f}")
    print(f"Precision: {prec:.4f}")
    print(f"Recall: {rec:.4f}")
    print(f"F1-score: {f1:.4f}")
    print(classification_report(y_test, y_pred, digits=4))

    cm = confusion_matrix(y_test, y_pred, labels=['spam', 'ham'])
    fig, ax = plt.subplots(figsize=(5, 4))
    im = ax.imshow(cm, interpolation='nearest')
    ax.set_title(f'Confusion Matrix - {model_name}')
    ax.set_xticks([0, 1])
    ax.set_yticks([0, 1])
    ax.set_xticklabels(['spam', 'ham'])
    ax.set_yticklabels(['spam', 'ham'])
    for i in range(cm.shape[0]):
        for j in range(cm.shape[1]):
            ax.text(j, i, cm[i, j], ha='center', va='center')

    plt.xlabel('Predicted')
    plt.ylabel('Actual')
    plt.tight_layout()
    plot_path = f"{plot_prefix}_{model_name.replace(' ', '_')}_cm.png"
    fig.savefig(plot_path)
    plt.close(fig)
    print(f'Confusion matrix saved to {plot_path}')

    return {'accuracy': acc, 'precision': prec, 'recall': rec, 'f1': f1}

```

```

def train_and_evaluate(df: pd.DataFrame, output_dir: str = './output') -> dict:
    os.makedirs(output_dir, exist_ok=True)
    X = df['text_clean'].values
    y = df['label'].values
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)

    vect = build_vectorizer()
    X_train_tfidf = vect.fit_transform(X_train)
    X_test_tfidf = vect.transform(X_test)

    models = {
        'MultinomialNB': MultinomialNB(),
        'LogisticRegression': LogisticRegression(max_iter=1000),
        'SVM_Linear': LinearSVC(max_iter=10000),
        'RandomForest': RandomForestClassifier(n_estimators=200),
        'KNN': KNeighborsClassifier(n_neighbors=5)
    }

    results = {}
    best_model = None
    best_score = -1
    best_name = None

    for name, clf in models.items():
        print(f'\nTraining {name} ...')
        clf.fit(X_train_tfidf, y_train)
        pipeline = Pipeline([('vect', vect), ('clf', clf)])
        metrics = evaluate_model(pipeline, X_test, y_test, model_name=name,
                                plot_prefix=os.path.join(output_dir, 'cm'))
        results[name] = metrics

        if metrics['f1'] > best_score:
            best_score = metrics['f1']
            best_model = pipeline
            best_name = name

    print(f'\nBest model: {best_name} with F1 = {best_score:.4f}')
    model_path = os.path.join(output_dir, 'model.joblib')
    joblib.dump(best_model, model_path)
    print(f'Saved best model pipeline to {model_path}')

    return {'results': results, 'best_model_name': best_name, 'model_path': model_path}

def create_flask_app(model_path: str):
    from flask import Flask, request, jsonify
    app = Flask('spam-detector')

```

```

if not os.path.exists(model_path):
    raise ValueError(f'Model file not found: {model_path}')

model = joblib.load(model_path)

@app.route('/predict', methods=['POST'])
def predict():
    data = request.get_json(force=True)
    text = data.get('text', '')
    text_clean = basic_clean(text)
    pred = model.predict([text_clean])[0]

    prob = None
    try:
        if hasattr(model.named_steps['clf'], 'predict_proba'):
            prob = model.named_steps['clf'].predict_proba(
                model.named_steps['vect'].transform([text_clean])
            )[0].tolist()
    except Exception:
        prob = None

    return jsonify({'prediction': pred, 'probability': prob})

@app.route('/', methods=['GET'])
def index():
    return "Spam Detection API. POST JSON {'text': '<email body>'} to /predict"

return app

def main(args):
    if args.mode == 'train':
        df = prepare_data(args.data_dir)
        res = train_and_evaluate(df, output_dir=args.output_dir)
        print('\nTraining complete. Summary:')
        for name, metrics in res['results'].items():
            print(f'{name}: F1={metrics['f1']:.4f}, Acc={metrics['accuracy']:.4f}')
        print(f'Model saved to: {res['model_path']}')

    elif args.mode == 'serve':
        if not args.model_path:
            raise ValueError('Please provide --model-path to a trained model.joblib')
        app = create_flask_app(args.model_path)
        app.run(host='0.0.0.0', port=args.port, debug=False)
    else:
        raise ValueError('Unknown mode. Use train or serve')

if __name__ == '__main__':
    parser = argparse.ArgumentParser(description='Email Spam Detection - Train or Serve')

```

```

parser.add_argument('--mode', choices=['train', 'serve'], required=True,
                    help='train: train model; serve: start API server')
parser.add_argument('--data-dir', default='./data', help='Data directory (CSV or folders)')
parser.add_argument('--output-dir', default='./output', help='Directory to save outputs')
parser.add_argument('--model-path', default='./output/model.joblib', help='Path to trained
model (for serve mode)')
parser.add_argument('--port', type=int, default=5000, help='Port for Flask server')
args = parser.parse_args()

# Ensure nltk stopwords available (optional)
if nltk is not None:
    try:
        _ = stopwords.words('english')
    except Exception:
        print('NLTK stopwords not found. Downloading...')
        nltk.download('stopwords')

```

```
main(args)
```

7. Results

The model achieved an accuracy of approximately 98.3% using the Linear SVM classifier, which performed best among all tested models. It effectively distinguishes between spam and legitimate (ham) emails by analyzing textual patterns and word frequencies using the TF-IDF feature extraction technique.


```
PS C:\Users\Ash\Projects\SpamDetector> python email_spam_detection.py --mode train --data-dir ./data
```

```
Loading dataset from ./data/emails.csv
```

```
Loaded 5171 records
```

```
Training MultinomialNB ...
```

```
== MultinomialNB ==
```

```
Accuracy: 0.9723
```

```
Precision: 0.9667
```

```
Recall: 0.9540
```

```
F1-score: 0.9603
```

	precision	recall	f1-score	support
ham	0.9783	0.9865	0.9824	846
spam	0.9667	0.9540	0.9603	279
accuracy			0.9723	1125
macro avg	0.9725	0.9703	0.9714	1125
weighted avg	0.9723	0.9723	0.9723	1125

```
Training LogisticRegression ...
```

```
== LogisticRegression ==
```

```
Accuracy: 0.9811
```

```
Precision: 0.9724
```

```
Recall: 0.9670
```

```
F1-score: 0.9697
```

	precision	recall	f1-score	support
ham	0.9864	0.9894	0.9879	846
spam	0.9724	0.9670	0.9697	279
accuracy			0.9811	1125
macro avg	0.9794	0.9782	0.9788	1125
weighted avg	0.9811	0.9811	0.9811	1125

Confusion matrix saved to ./output/cm_LogisticRegression_cm.png

Training SVM_Linear ...

== SVM_Linear ==

Accuracy: 0.9835

Precision: 0.9743

Recall: 0.9715

F1-score: 0.9729

	precision	recall	f1-score	support
ham	0.9880	0.9910	0.9895	846
spam	0.9743	0.9715	0.9729	279
accuracy			0.9835	1125
macro avg	0.9811	0.9812	0.9812	1125
weighted avg	0.9835	0.9835	0.9835	1125

Confusion matrix saved to ./output/cm_SVM_Linear_cm.png

Training RandomForest ...

== RandomForest ==

Accuracy: 0.9787

Precision: 0.9620

Recall: 0.9650

F1-score: 0.9634

	precision	recall	f1-score	support
ham	0.9830	0.9845	0.9837	846
spam	0.9620	0.9650	0.9634	279
accuracy			0.9787	1125
macro avg	0.9725	0.9747	0.9736	1125
weighted avg	0.9787	0.9787	0.9787	1125

Confusion matrix saved to ./output/cm_RandomForest_cm.png

Training KNN ...

== KNN ==

Accuracy: 0.9568

Precision: 0.9400

Recall: 0.9400

F1-score: 0.9391

	precision	recall	f1-score	support
ham	0.9720	0.9710	0.9715	846
spam	0.9400	0.9400	0.9391	279
accuracy			0.9568	1125
macro avg	0.9560	0.9555	0.9553	1125
weighted avg	0.9568	0.9568	0.9568	1125

Confusion matrix saved to ./output/cm_KNN_cm.png

Best model: SVM_Linear with F1 = 0.9729

Saved best model pipeline to ./output/model.joblib

Training complete. Summary:

MultinomialNB: F1=0.9603, Acc=0.9723

LogisticRegression: F1=0.9697, Acc=0.9811

SVM_Linear: F1=0.9729, Acc=0.9835

RandomForest: F1=0.9634, Acc=0.9787

KNN: F1=0.9391, Acc=0.9568

Model saved to: ./output/model.joblib

8. Conclusion

The Email Spam Detection System successfully demonstrates the application of machine learning in filtering unwanted emails. By preprocessing email text and extracting TF-IDF features, the system was able to train and compare multiple classification algorithms. Among them, the **Linear SVM model** achieved the highest performance with an accuracy of around **98.3%**, effectively distinguishing spam from legitimate messages. This approach provides a reliable, automated, and scalable solution to enhance email security and reduce user exposure to phishing and advertising spam..